

# **Hospital Readmission**

## **Prediction Using**

## **Machine Learning**

presented by:

VEJJU KRISHNA SAI RUTHVIK  
ATKURI SAI SURYA  
BASSETTI ADINARAYANA  
KAKUMANU GEETHA MADHAV

# **1. INTRODUCTION :**

## **a) Overview:**

If a hospital has multiple readmissions, it means that the hospital needs to work on the quality of services it is providing with respect to the health and wellness of its patients. Being able to predict whether a person will be readmitted to the hospital within 30 days or not, will be of great help to the hospital in developing an idea of the incoming number of repeated patients which in turn helps to provide better services for patients with increased risk of disease.....

One patient population that is at increased risk of hospitalization and readmission is diabetes. Diabetes is a medical condition that affects approximately 1 in 10 patients in the United States. So in this project, we will be focusing on hospital readmission prediction for patients who are having diabetes.

This study used the Health Facts database (Cerner Corporation, Kansas City, MO), a national data warehouse that collects comprehensive clinical records across hospitals throughout the United States. The Health Facts data we used was an extract representing 10 years (1999–2008) of clinical care at 130 hospitals and integrated delivery networks throughout the United States.

## **b) Purpose:**

The main purpose of this project is to predict whether a person who is suffering from diabetes and consulting a specific hospital will be readmitted or not, based on multiple factors.

We will be using classification algorithms such as Logistic Regression, KNN, Decision tree, Random Forest, AdaBoost, and GradientBoost. We will train and test the data with these algorithms. From this, the best model is selected and saved in pkl format. We will also be

deploying our model locally using Flask.

## **2. LITERATURE SURVEY:**

### **Some of the existing problem:**

H2O.ai:

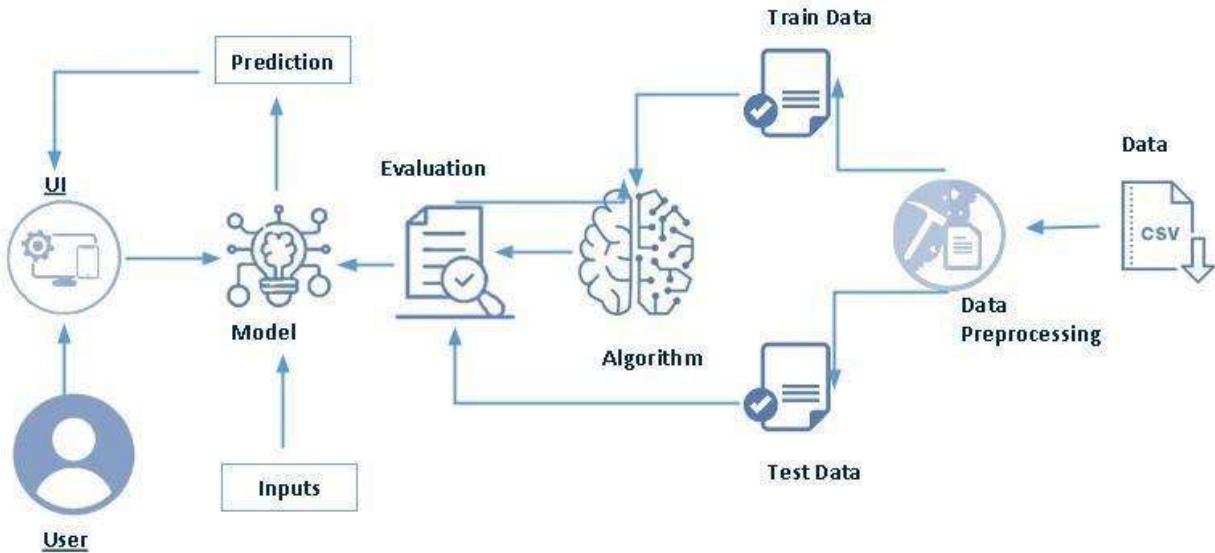
The mission at H2O.ai is to democratize AI for all so that more people across industries can use the power of AI to solve business and social challenges. The healthcare industry is a key focus for the company with an initiative to help develop AI healthcare solutions including dedicated, experienced resources for customers, driving healthcare AI events and meetups for healthcare professionals, and membership in Health IT Now, the leading coalition of patient groups, provider organizations, employers, insurers, and other stakeholders. H2O.ai is already working with top healthcare companies including Change Healthcare, Armada Health, Kaiser Permanente, and HCA, and its products include industry leading features for machine learning interpretability required by the healthcare industry for compliance purposes.

### **Proposed Solution:**

The main purpose of this project is to predict whether a person who is suffering from diabetes and consulting a specific hospital will be readmitted or not, based on multiple factors.

## **3. THEORITICAL ANALYSIS:**

### **a) Block Diagram:**



## b) Hard/Software Designing:

Operating System	Windows, Mac, Linux
CPU (for training)	Multi Core Processors (i3 or above/equivalent)
GPU (for training)	NVIDIA AI Capable / Google's TPU

Software Requirements:

Python	v3.10.0 or Above
Python Packages	flask, tensorflow, opencv-python, keras, numpy, pandas, virtualenv, pillow
Web Browser	Mozilla Firefox, Google Chrome or any modern web browser

## 4. EXPERIMENTAL INVESTIGATIONS:

### Training the train dataset:

The screenshot shows a Jupyter Notebook interface running on a local host. The notebook file is titled "Hospital Readmission.ipynb". The code cell contains Python code for importing various machine learning classifiers from the scikit-learn library and defining a function to test the models. The output cell shows the execution of the code and the resulting accuracy scores for each model.

```
[101]: from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
from xgboost.sklearn import XGBClassifier
from sklearn import metrics
from sklearn.metrics import accuracy_score, roc_curve, confusion_matrix, classification_report, auc

[135]: model_dict = {}

model_dict['Logistic regression'] = LogisticRegression(solver='liblinear', random_state=20)
model_dict['KNN Classifier'] = KNeighborsClassifier()
model_dict['Decision Tree Classifier'] = DecisionTreeClassifier(random_state=20)
model_dict['Random Forest Classifier'] = RandomForestClassifier(random_state=20)
model_dict['AdaBoost Classifier'] = AdaBoostClassifier(random_state=20)
model_dict['Gradient Boosting Classifier'] = GradientBoostingClassifier(random_state=20)

[136]: def model_test(X_train, X_test, y_train, y_test, model, model_name):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print('===== {} ====='.format(model_name))
    print('Score is : {}'.format(accuracy))
    print()

[137]: for model_name, model in model_dict.items():
    model_test(X_train, X_val, y_train, y_val, model, model_name)

===== Logistic regression =====
Score is : 0.5861663108337439
===== KNN Classifier =====
```

localhost:8888/lab/tree/Hospital%20Readmission%20Prediction/notebook/Hospital%20Readmission.ipynb

```
[137]: for model_name, model in model_dict.items():
    model_test(X_train, X_val, y_train, y_val, model, model_name)

=====logistic regression=====
Score is : 0.5861663108337439

=====KNN Classifier=====
Score is : 0.7962719658400395

=====Decision Tree Classifier=====
Score is : 0.8781956533639897

=====Random Forest Classifier=====
Score is : 0.917309903103958

=====AdaBoost Classifier=====
Score is : 0.818333607950962

=====Gradient Boosting Classifier=====
Score is : 0.8846280177369026
```

```
[138]: for model_name, model in model_dict.items():
    p = model.predict(X_test)
    print('Testing accuracy of ',model_name,'%',accuracy_score(y_test,p))

Testing accuracy of Logistic regression = 0.66837775356034811
Testing accuracy of KNN Classifier = 0.6339067174326897
Testing accuracy of Decision Tree Classifier = 0.7998368234974164
Testing accuracy of Random Forest Classifier = 0.8569146042969812
Testing accuracy of AdaBoost Classifier = 0.85878433955752
Testing accuracy of Gradient Boosting Classifier = 0.8802352461245581
```

```
[144]: cf_matrix = confusion_matrix(y_val, pred_rfc)
sb.heatmap(cf_matrix/np.sum(cf_matrix), annot=True, fmt='%.2%')
```

Mode: Command Ln 1, Col 1 Hospital Readmission.ipynb

localhost:8888/lab/tree/Hospital%20Readmission%20Prediction/notebook/Hospital%20Readmission.ipynb

```
[144]: cf_matrix = confusion_matrix(y_val, pred_rfc)
sb.heatmap(cf_matrix/np.sum(cf_matrix), annot=True, fmt='%.2%')
plt.title("Confusion Matrix of Random Forest Classifier")
plt.xlabel("Predicted")
plt.ylabel("Actual")
```

```
[146]: fpr_rfc, tpr_rfc, thresholds_rfc = roc_curve(y_val, pred_rfc)
roc_auc_rfc = metrics.auc(fpr_rfc, tpr_rfc)
plt.plot(fpr_rfc, tpr_rfc, color="orange", label="ROC curve (area = %0.2f)" % roc_auc_rfc)
plt.plot([0, 1], [0, 1], color="blue", linestyle="--")
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.title('ROC Curve')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.legend(loc="lower right")
```

Mode: Command Ln 1, Col 1 Hospital Readmission.ipynb

localhost:8888/lab/tree/Hospital%20Readmission%20Prediction/notebook/Hospital%20Readmission.ipynb

```
[105]: rfc = RandomForestClassifier(random_state=28)
rfc.fit(Xtrain,y_train)
pred_rfc = rfc.predict(X_val)
print("Training Accuracy of Random Forest:",accuracy_score(y_val,pred_rfc))

Training Accuracy of Random Forest: 0.938632506706082

[108]: cols = list(data.columns)
cols.remove('readmitted')

[109]: importance = rfc.feature_importances_
for i,v in enumerate(importance):
    print("Feature: %d, Score: %.5f" % (i,v))
print(cols)
plt.figure(figsize=(15,10))
plt.bar(cols[x] for x in range(len(importance))), importance
plt.xticks(rotation=90)
plt.show()

Feature: 0, Score: 0.02612
Feature: 1, Score: 0.02665
Feature: 2, Score: 0.02879
Feature: 3, Score: 0.02301
Feature: 4, Score: 0.07155
Feature: 5, Score: 0.08739
Feature: 6, Score: 0.05176
Feature: 7, Score: 0.02156
Feature: 8, Score: 0.01079
Feature: 9, Score: 0.07860
Feature: 10, Score: 0.01159
Feature: 11, Score: 0.03947
Feature: 12, Score: 0.03854
Feature: 13, Score: 0.03565
Feature: 14, Score: 0.03560
Feature: 15, Score: 0.01886
```

Simple 0 1 Python 3 (ipykernel) | Idle Mode: Command Ln 1, Col 1 Hospital Readmission.ipynb

localhost:8888/lab/tree/Hospital%20Readmission%20Prediction/notebook/Hospital%20Readmission.ipynb

```
[146]: (array([0., 1.]), array([0., 1.]), array([2, 1, 0], dtype=int64))

[147]: print(classification_report(y_val,pred_rfc))
      precision    recall  f1-score   support

          0       0.89      1.00      0.94     18338
          1       1.00      0.88      0.93     18196

   accuracy                           0.94      36534
  macro avg       0.95      0.94      0.94      36534
weighted avg       0.94      0.94      0.94      36534
```

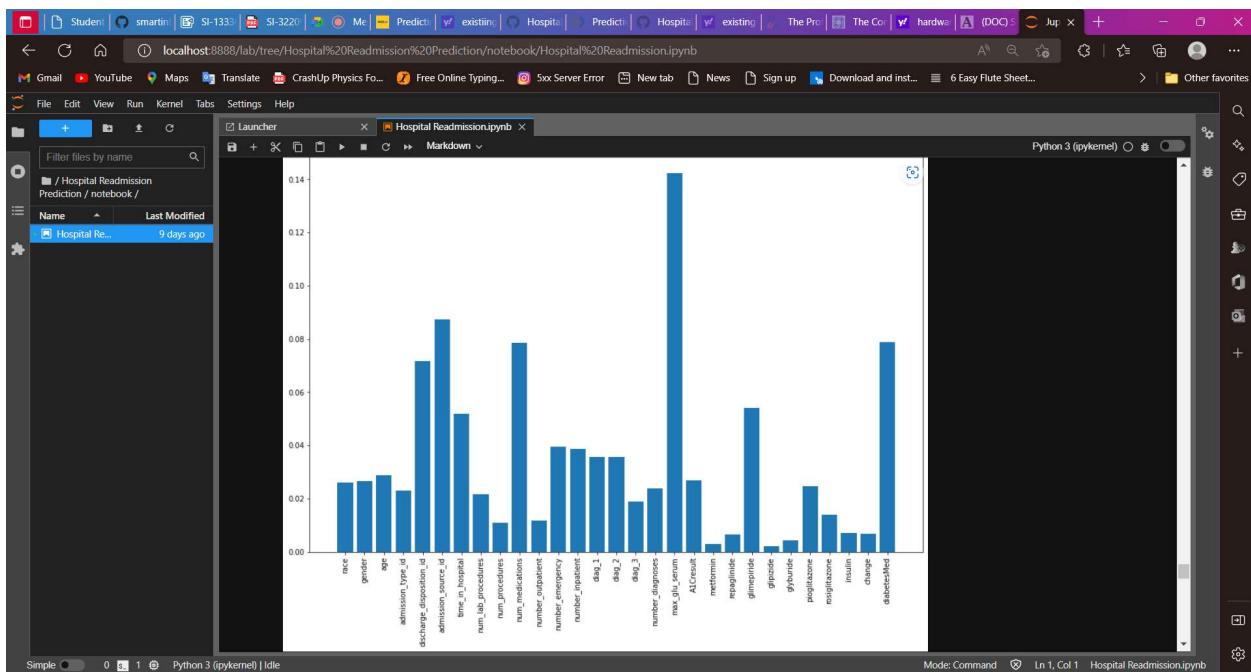
[106]: rfc = RandomForestClassifier(random\_state=28)

Simple 0 1 Python 3 (ipykernel) | Idle Mode: Command Ln 1, Col 1 Hospital Readmission.ipynb

localhost:8888/lab/tree/Hospital%20Readmission%20Prediction/notebook/Hospital%20Readmission.ipynb

```
plt.show()

Feature: 0, Score: 0.02612
Feature: 1, Score: 0.02665
Feature: 2, Score: 0.02879
Feature: 3, Score: 0.02701
Feature: 4, Score: 0.07155
Feature: 5, Score: 0.028729
Feature: 6, Score: 0.05176
Feature: 7, Score: 0.02156
Feature: 8, Score: 0.01079
Feature: 9, Score: 0.07869
Feature: 10, Score: 0.01159
Feature: 11, Score: 0.03947
Feature: 12, Score: 0.03854
Feature: 13, Score: 0.03565
Feature: 14, Score: 0.03569
Feature: 15, Score: 0.01888
Feature: 16, Score: 0.02385
Feature: 17, Score: 0.14236
Feature: 18, Score: 0.02681
Feature: 19, Score: 0.00295
Feature: 20, Score: 0.00651
Feature: 21, Score: 0.05393
Feature: 22, Score: 0.00209
Feature: 23, Score: 0.00441
Feature: 24, Score: 0.02454
Feature: 25, Score: 0.01401
Feature: 26, Score: 0.00689
Feature: 27, Score: 0.00687
Feature: 28, Score: 0.07883
['race', 'gender', 'age', 'admission_type_id', 'discharge_disposition_id', 'admission_source_id', 'time_in_hospital', 'num_lab_procedures', 'num_procedures', 'num_medications', 'number_outpatient', 'number_inpatient', 'diag_1', 'diag_2', 'diag_3', 'number_diagnoses', 'max_glu_serum', 'AlCresult', 'metformin', 'repaglinide', 'glimepiride', 'glyburide', 'pioglitazone', 'rosiglitazone', 'insulin', 'change', 'diabetesMed']
```



localhost:8888/lab/tree/Hospital%20Readmission%20Prediction/notebook/Hospital%20Readmission.ipynb

```
[217]: p = rfc.predict(X_test)
print('Testing Accuracy of Random Forest: ',accuracy_score(y_test,p))

Testing Accuracy of Random Forest:  0.85691466042969812

[183]: rfc.predict([[ 8.,  1.,  1.,  1.,  7.,  4., 24.,  0.,  4.,  1.,  0.]])
```

```
[183]: array([1], dtype=int64)
```

```
[188]: imp_cols = []
for i,v in enumerate(importance):
    if (v>0.03):
        imp_cols.append(cols[i])
```

```
[111]: imp_cols
```

```
[111]: ['discharge_disposition_id',
       'admission_source_id',
       'time_in_hospital',
       'num_medications',
       'number_emergency',
       'number_inpatient',
       'diag_1',
       'diag_2',
       'max_glu_serum',
       'glimepiride',
       'diabetesMed']
```

```
[120]: final_data = data[imp_cols['readmitted']].copy()
```

```
[121]: final_data.shape
```

```
[121]: (98052, 12)
```

```
[123]: X1 = final_data.drop('readmitted', axis = 1)
y1 = final_data['readmitted']
```

```
[123]: ct = ColumnTransformer([('oe',OrdinalEncoder()),['diag_1','diag_2','diabetesMed']])
```

```
[123]: ct.fit_transform(X1)
```

```
[124]: X1 = ct.transform(X1)
```

```
[125]: X1.shape, y1.shape
```

```
[125]: ((98052, 11), (98052,))
```

```
[126]: X_train, X_test, y_train, y_test = train_test_split(X1,y1,
                                                       stratify=y1,
                                                       test_size=0.3, random_state=20)
```

```
[127]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
[127]: ((68636, 11), (29416, 11), (68636,), (29416,))
```

```
[304]: X_test[20:40]
```

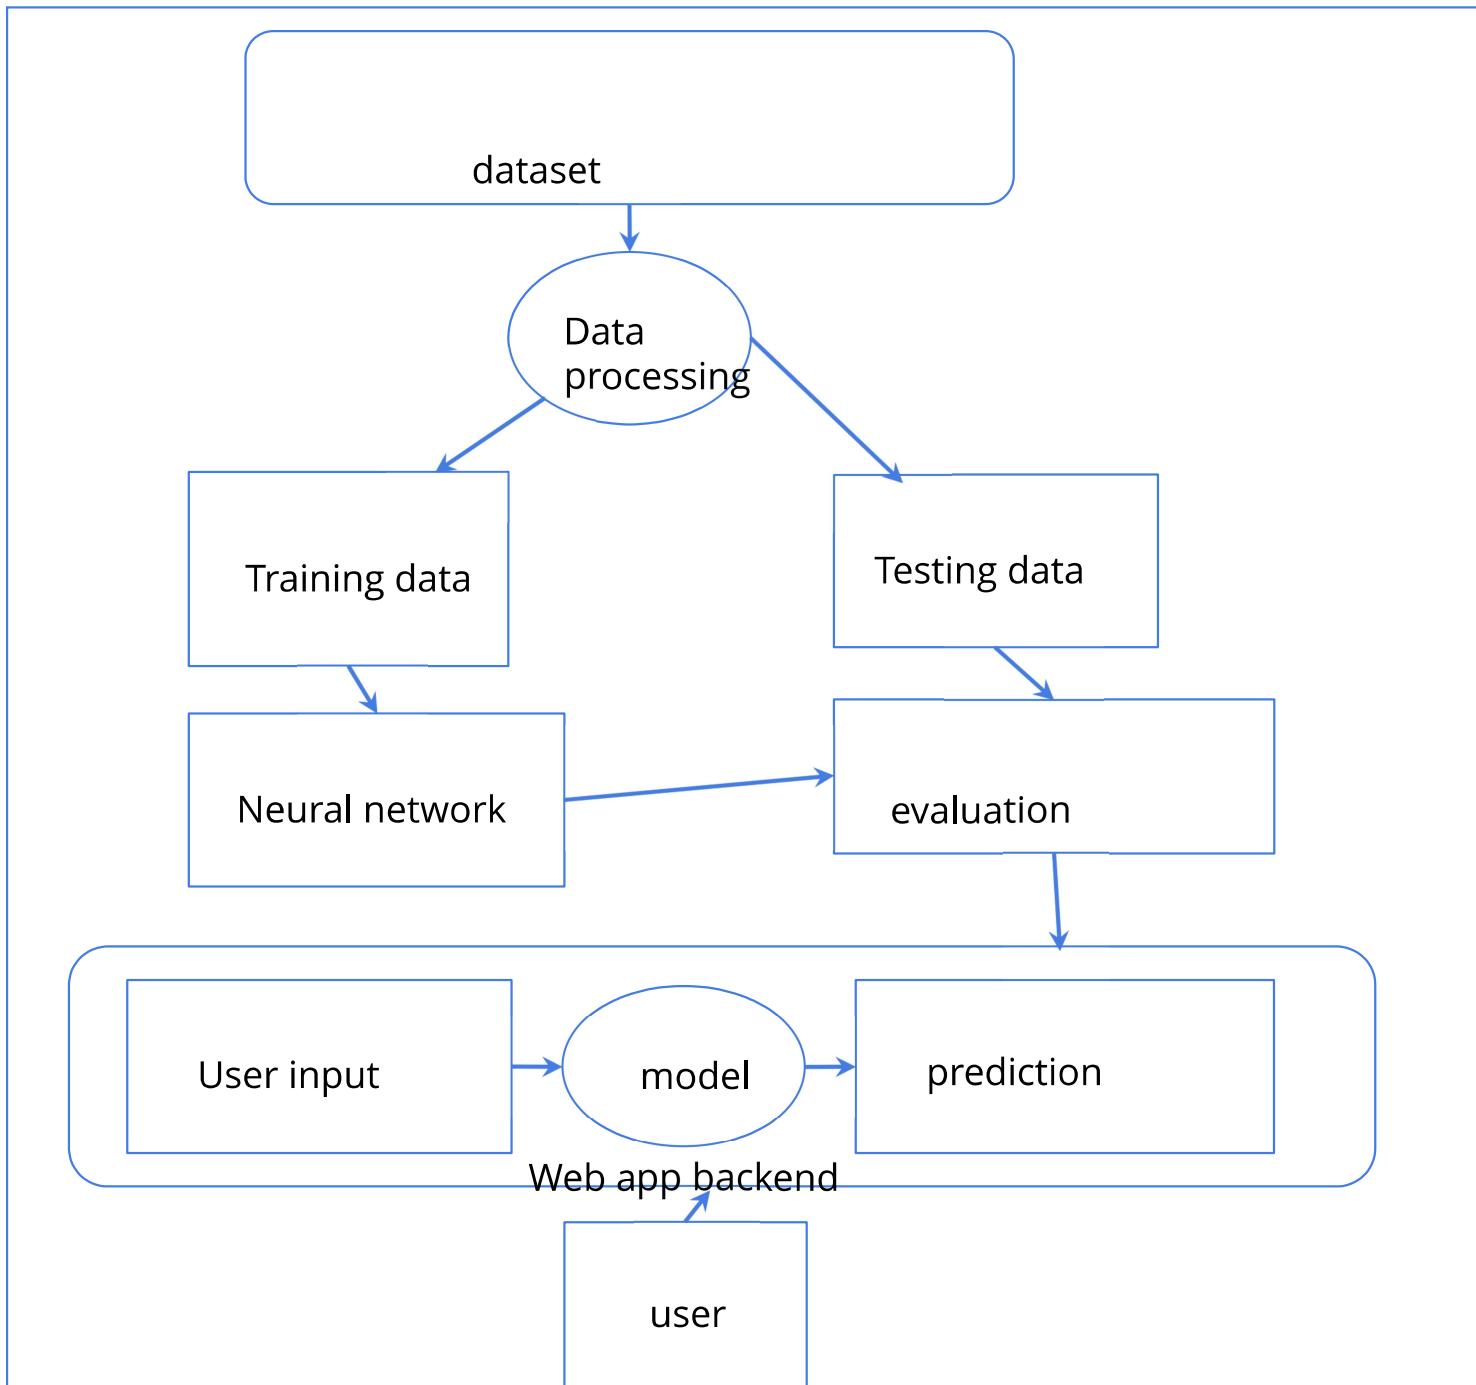
```
[304]: array([[ 7.,  3.,  1.,  2.,  7.,  9., 22.,  0.,  0., 99.,  0.],
   [ 2.,  8.,  1.,  1.,  7.,  8., 19.,  0.,  1., 99.,  0.],
   [ 0.,  0.,  1.,  1.,  7.,  2., 26.,  0.,  0., 99.,  0.],
   [ 8.,  1.,  1.,  1.,  7.,  4., 24.,  0.,  4., 1.,  0.],
   [ 0.,  0.,  1.,  1.,  7.,  3., 18.,  0.,  2., 99.,  0.],
   [ 4.,  0.,  1.,  1.,  7.,  6., 14.,  0.,  0., 99.,  0.],
   [ 8.,  2.,  1.,  1.,  7.,  4., 14.,  0.,  0., 99.,  0.],
   [ 3.,  0.,  1.,  1.,  7.,  4., 14.,  0.,  1., 99.,  0.],
   [ 8.,  8.,  0.,  1.,  1.,  5., 24.,  0.,  0., 99.,  0.],
   [ 0.,  0.,  0.,  1.,  1.,  1., 13.,  0.,  0., 99.,  0.],
   [ 0.,  0.,  1.,  18.,  1.,  4., 36.,  0.,  0., 99.,  0.],
   [ 7.,  7.,  1., 11.,  7.,  2., 11.,  0.,  0., 99.,  0.],
   [ 8.,  8.,  0.,  2.,  7.,  6., 14.,  0.,  0., 99.,  0.],
   [ 7.,  8.,  1.,  2.,  1.,  7., 26.,  0.,  0., 99.,  0.],
   [ 7.,  8.,  0.,  2.,  1.,  7., 26.,  0.,  0., 99.,  0.],
   [ 2.,  0.,  1.,  1.,  2., 12., 21.,  0.,  4., 99.,  0.],
   [ 2.,  7.,  0.,  2.,  7.,  5., 18.,  1.,  3., 99.,  0.],
   [ 7.,  1.,  1.,  1.,  9.,  9., 0.,  2., 99.,  0.],
```

The screenshot shows a Jupyter Notebook interface running in a browser window. The title bar indicates the URL is `localhost:8888/lab/tree/Hospital%20Readmission%20Prediction/notebook/Hospital%20Readmission.ipynb`. The notebook tab is titled "Hospital Readmission.ipynb". The code cell [305] contains the following Python code:

```
[305]: y_test[20:40]
11137 0
72049 0
68519 1
40887 0
26424 0
45104 0
71384 1
96221 0
38217 0
16568 0
11887 0
47468 0
50327 0
11843 0
44765 1
89439 0
68835 0
71002 0
30218 0
Name: readmitted, dtype: int64
```

The output of the code cell is displayed below the code, showing the first 20 rows of the `y_test` variable. The data consists of integer values (0 or 1) representing readmission status.

## 5. FLOWCHART:

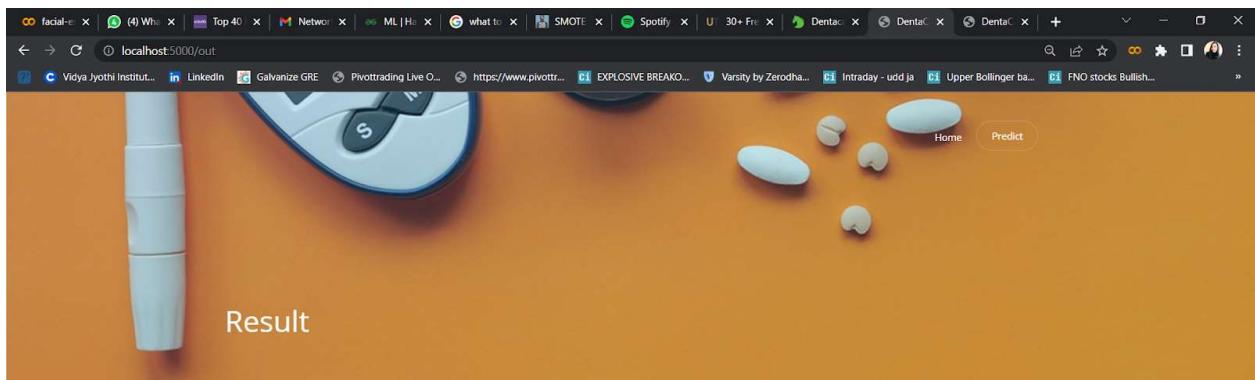


## 6. RESULT:

The problem associated is a binary classification. Since the target variable in the dataset is quite imbalanced, SMOTE algorithm was used to oversample the minority class.

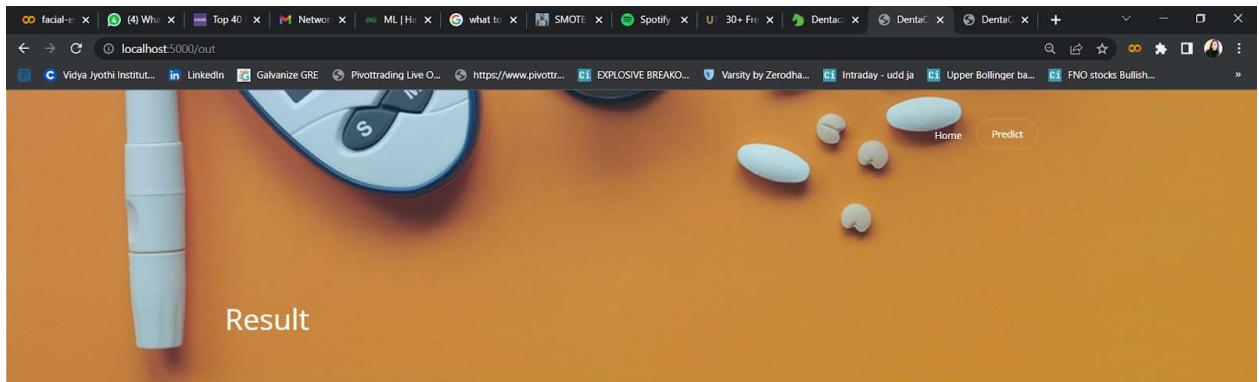
A number of models were used such as logistic regression, KNN Classifier, Decision Tree Classifier, Random Tree Classifier, Adaboost Classifier etc. We will train and test the data with these algorithms. From this, the best model is selected and saved in pkl format. We will also be deploying our model locally using Flask. Out of all the models, random forest classifier is found to perform the best.

The accuracy for Random Forest Classifier is 91.73 percent.



This patient will be readmitted





This patient will not be readmitted



## 7. ADVANTAGES AND DISADVANTAGES:

### Advantages:

1. The hospital will be able to be ready for any number of patients that will be readmitted in the future.

### Disadvantage:

1. The accuracy is only 91.73 percentage.
2. It will not be able to predict 100% accurately.

## **Applications:**

1. It can be used by clinics related to hospitals
2. It can be used by diabetes hospitals
3. It can be used by diabetic patients

## **7 CONCLUSION**

In this work we adopted machine learning methods to identify high risk patients and evaluated different machine learning algorithms. Compared to the previous analysis, our study achieved high accuracy due to the sophisticated pre processing procedure. The XGBoost method is reported to be the best method for prediction of the readmission rate for diabetes patients.

We identified the most important factors as the time\_in\_hospital and number of inpatient, number of diagnosis, which appears to associate with the severity of the disease. Further studies could conduct more exploration when analyzing these factors individually.

## **Future Scope:**

Having a technology that can predict the number of people that will be readmitted in the future will help the hospitals in being ready for any number of patients that will be readmitted.

## **Bibliography:**

- ML Concepts

Supervised learning: <https://www.youtube.com/watch?v=QeKshry8pWQ>

Unsupervised learning:

[https://www.youtube.com/watch?v=D6gtZrsYi6c&feature=emb\\_logo](https://www.youtube.com/watch?v=D6gtZrsYi6c&feature=emb_logo)

Metrics: <https://www.youtube.com/watch?v=aWAnNHXIKww&t=1s>

Flask : [https://www.youtube.com/watch?v=lj4l\\_CvBnt0&feature=emb\\_logo](https://www.youtube.com/watch?v=lj4l_CvBnt0&feature=emb_logo)

## APPENDIX:

```

1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(X, y,
3                                                 stratify=y,
4                                                 test_size=0.3,random_state=20)

1 X_train.shape, X_test.shape, y_train.shape, y_test.shape
((68636, 29), (29416, 29), (68636,), (29416,))

1 from imblearn.over_sampling import SMOTE
1 sm = SMOTE(random_state = 20)
2 X1_res, y1_res = sm.fit_sample(X_train,y_train)

1 X1_res = np.array(X1_res)
1 y1_res = np.array(y1_res)

1 X_train, X_val, y_train, y_val = train_test_split(X1_res,y1_res,test_size=0.3, random_state=20)

1 X_train.shape, X_val.shape, y_train.shape, y_val.shape
((85246, 29), (36534, 29), (85246,), (36534,))
```

```

1 from imblearn.over_sampling import SMOTE

1 sm = SMOTE(random_state = 20)
2 X1_res, y1_res = sm.fit_sample(X_train,y_train)

1 X1_res = np.array(X1_res)

1 y1_res = np.array(y1_res)

1
2 X_train, X_val, y_train, y_val = train_test_split(X1_res,y1_res,test_size=0.3, random_state=20)

1 X_train.shape, X_val.shape, y_train.shape, y_val.shape
((85246, 29), (36534, 29), (85246,), (36534,))

1 from sklearn.linear_model import LogisticRegression
2 from sklearn.tree import DecisionTreeClassifier
3 from sklearn.neighbors import KNeighborsClassifier
4 from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
5 from xgboost.sklearn import XGBClassifier
6 from sklearn import metrics
7 from sklearn.metrics import accuracy_score, roc_curve, confusion_matrix, classification_report, auc

1 #Creating a dictionary of all models
2 model_dict = {}
3
4 model_dict['Logistic regression']= LogisticRegression(solver='liblinear',random_state=20)
5 model_dict['KNN Classifier'] = KNeighborsClassifier()
6 model_dict['Decision Tree Classifier'] = DecisionTreeClassifier(random_state=20)
7 model_dict['Random Forest Classifier'] = RandomForestClassifier(random_state=20)
8 model_dict['AdaBoost Classifier'] = AdaBoostClassifier(random_state=20)
9 model_dict['Gradient Boosting Classifier'] = GradientBoostingClassifier(random_state=20)
10 model_dict['XGB Classifier'] = XGBClassifier(random_state=20)

```

```
1 #function to print accuracy of all models
2 def model_test(X_train, X_test, y_train, y_test, model, model_name):
3     model.fit(X_train, y_train)
4     y_pred = model.predict(X_test)
5     accuracy = accuracy_score(y_test, y_pred)
6     print('-----{}-----'.format(model_name))
7     print('Score is : {}'.format(accuracy))
8
9     print()
```

```
1 for model_name, model in model_dict.items():
2     model_test(X_train, X_val, y_train, y_val, model, model_name)
```

```
=====Logistic regression=====
Score is : 0.6002080254009964
```

```
=====KNN Classifier=====
Score is : 0.7935621612744292
```

```
=====Decision Tree Classifier=====
Score is : 0.8801664203207971
```

```
=====Random Forest Classifier=====
Score is : 0.938632506706082
```

```
=====AdaBoost Classifier=====
Score is : 0.9139158044561231
```

```
=====Gradient Boosting Classifier=====
Score is : 0.9317621941205453
```

```
=====XGB Classifier=====
Score is : 0.9375376361745223
```

```
1 for model_name, model in model_dict.items():
2     p = model.predict(X_test)
3     print('Testing accuracy of ', model_name, '=', accuracy_score(y_test, p))
```

```
Testing accuracy of Logistic regression = 0.6484226271416916
```

```
Testing accuracy of KNN Classifier = 0.5848517813434866
```

```
Testing accuracy of Decision Tree Classifier = 0.7861028011966277
```

```
Testing accuracy of Random Forest Classifier = 0.8871702474843622
```

```
Testing accuracy of AdaBoost Classifier = 0.8744220832200164
```

```
Testing accuracy of Gradient Boosting Classifier = 0.8873742181125918
```

```
Testing accuracy of XGB Classifier = 0.8872382376937721
```

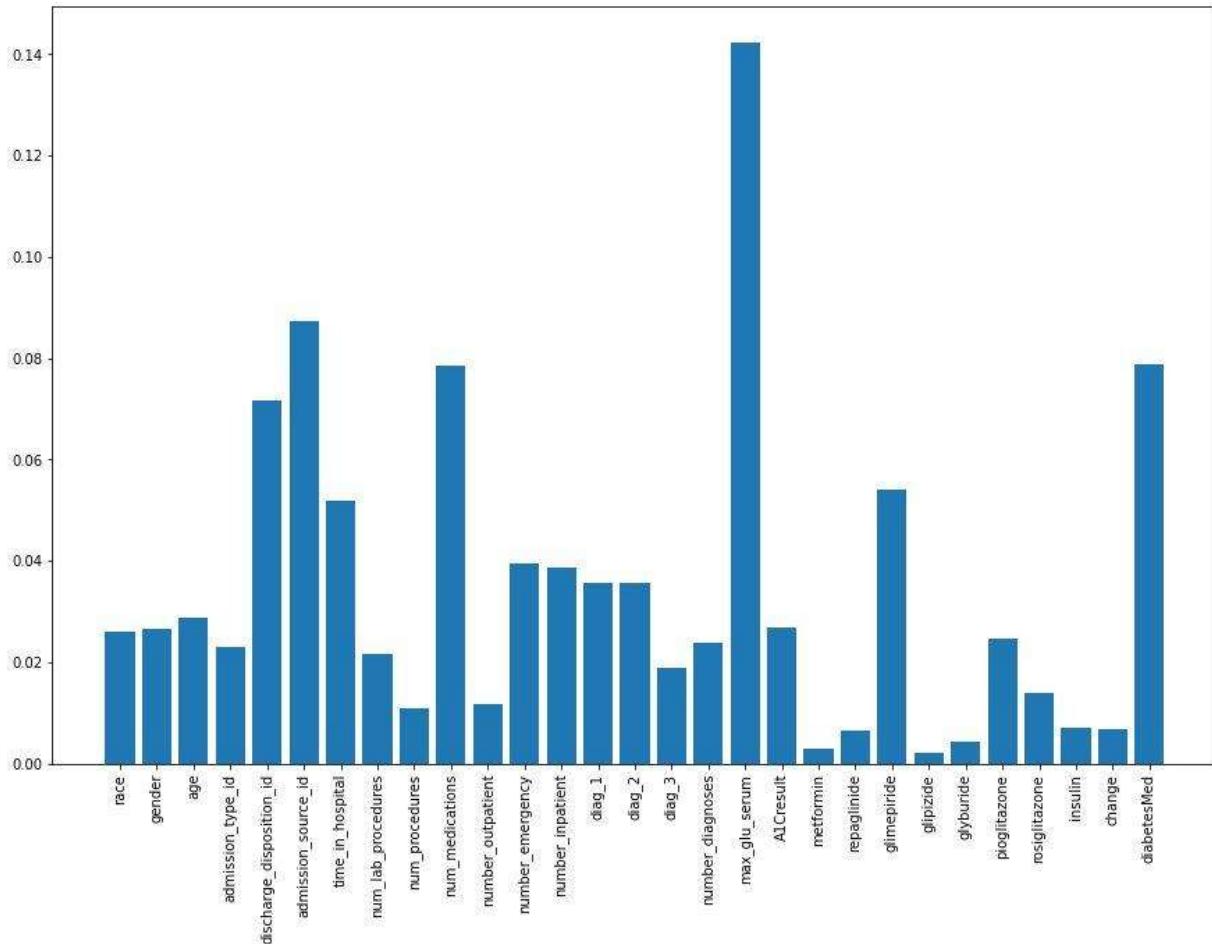
```
1 #Fitting data to Random Forest classifier
2 rfc = RandomForestClassifier(random_state=20)
3 rfc.fit(X_train, y_train)
4 pred_rfc = rfc.predict(X_val)
5 print('Training Accuracy of Random Forest=', accuracy_score(y_val, pred_rfc))
```

```
Training Accuracy of Random Forest= 0.938632506706082
```

```

1 importance = rfc.feature_importances_
2 # summarize feature importance
3 for i,v in enumerate(importance):
4     print('Feature: %0d, Score: %.5f' % (i,v))
5 # plot feature importance
6 print(cols)
7 plt.figure(figsize=(15,10))
8 plt.bar([cols[x] for x in range(len(importance))], importance)
9 plt.xticks(rotation=90)
10 plt.show()

```



```
1 imp_cols = []
2 for i,v in enumerate(importance):
3     if (v>0.035):
4         imp_cols.append(cols[i])
```

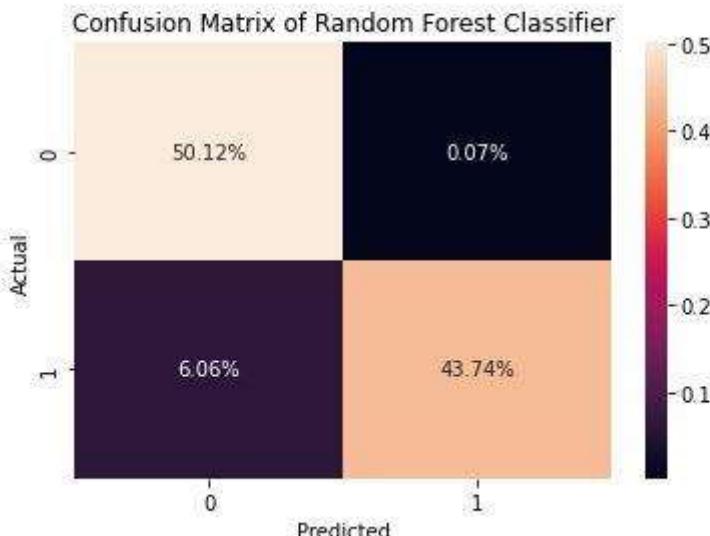
```
1 imp_cols
```

```
['discharge_disposition_id',
 'admission_source_id',
 'time_in_hospital',
 'num_medications',
 'number_emergency',
 'number_inpatient',
 'diag_1',
 'diag_2',
 'max_glu_serum',
 'glimepiride',
 'diabetesMed']
```

Confusion matrix:

```
1 #Plotting confusion matrix
2 cf_matrix = confusion_matrix(y_val, pred_rfc)
3 sb.heatmap(cf_matrix/np.sum(cf_matrix), annot=True, fmt='.2%')
4 plt.title('Confusion Matrix of Random Forest Classifier')
5 plt.xlabel('Predicted')
6 plt.ylabel('Actual')
```

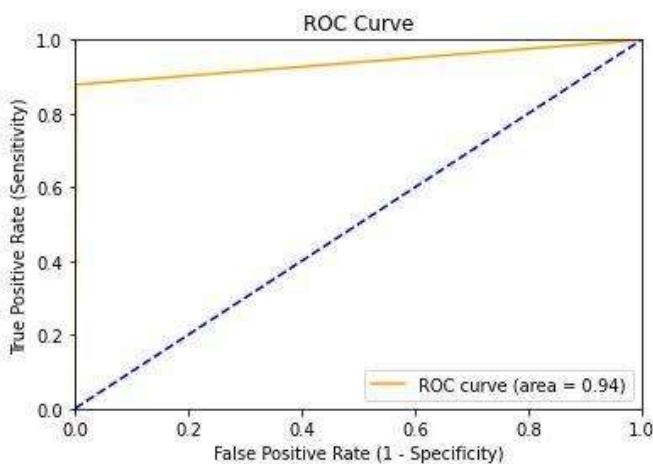
```
Text(33.0, 0.5, 'Actual')
```



```

1 #Plotting ROC curve
2 fpr_rfc, tpr_rfc, thresholds_rfc = roc_curve(y_val, pred_rfc)
3 roc_auc_rfc = metrics.auc(fpr_rfc, tpr_rfc)
4 plt.plot(fpr_rfc, tpr_rfc, color='orange', label='ROC curve (area = %0.2f)' % roc_auc_rfc)
5 plt.plot([0, 1], [0, 1], color='blue', linestyle='--')
6 plt.xlim([0.0, 1.0])
7 plt.ylim([0.0, 1.0])
8 plt.title('ROC Curve')
9 plt.xlabel('False Positive Rate (1 - Specificity)')
10 plt.ylabel('True Positive Rate (Sensitivity)')
11 plt.legend(loc="lower right")
12 plt.show()
13 roc_curve(y_val, pred_rfc)

```



```
1 print(classification_report(y_val,pred_rfc))
```

	precision	recall	f1-score	support
0	0.89	1.00	0.94	18338
1	1.00	0.88	0.93	18196
accuracy			0.94	36534
macro avg	0.95	0.94	0.94	36534
weighted avg	0.94	0.94	0.94	36534

```

1 import pickle
2 pickle.dump(rfc,open('model.pkl','wb'))

```

Web application in Flask:

```
from flask import Flask, render_template, request
import pickle, joblib
import pandas as pd

app = Flask(__name__)

model = pickle.load(open("model.pkl", "rb"))
ct = joblib.load('feature_values')

@app.route('/')
def home():
    return render_template("home.html")

@app.route('/pred')
def predict():
    return render_template("index.html")
```

```

@app.route('/out', methods =["POST"])
def output():
    discharge_disposition_id = request.form["discharge_disposition_id"]
    admission_source_id = request.form["admission_source_id"]
    time_in_hospital = request.form["time_in_hospital"]
    num_medications = request.form["num_medications"]
    number_emergency = request.form["number_emergency"]
    number_inpatient = request.form["number_inpatient"]
    diag_1 = request.form["diag_1"]
    diag_2 = request.form["diag_2"]
    max_glu_serum = request.form["max_glu_serum"]
    glimepiride = request.form["glimepiride"]
    diabetesMed = request.form["diabetesMed"]

    if max_glu_serum == '>200' or max_glu_serum=='>300':
        max_glu_serum=1
    elif max_glu_serum=='Norm':
        max_glu_serum=0
    else:
        max_glu_serum=-99

    if glimepiride == 'No':
        glimepiride = 0
    else:
        glimepiride=1

    data = [[discharge_disposition_id,admission_source_id,time_in_hospital,
             num_medications, number_emergency, number_inpatient,diag_1, diag_2,
             max_glu_serum, glimepiride, diabetesMed]]

    feature_cols = ['discharge_disposition_id', 'admission_source_id', 'time_in_hospital',
                    'num_medications', 'number_emergency', 'number_inpatient',
                    'diag_1', 'diag_2','max_glu_serum', 'glimepiride', 'diabetesMed']

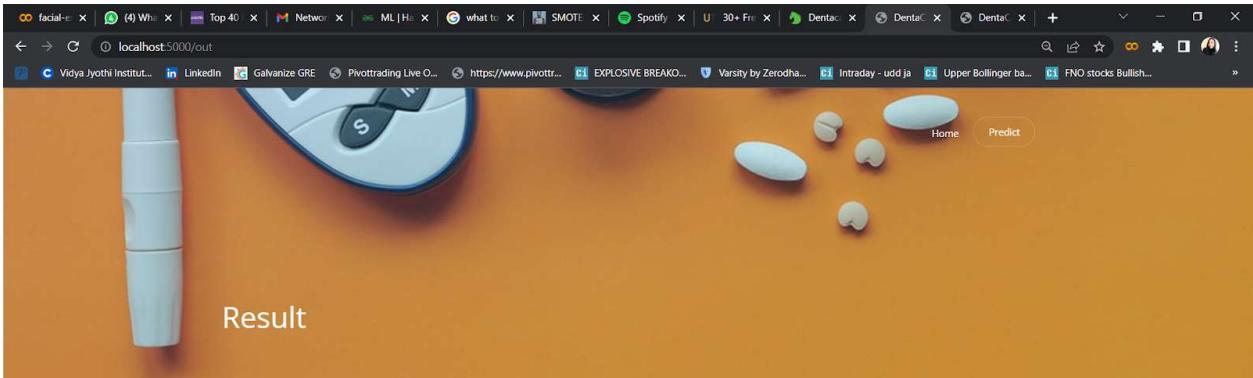
    pred = model.predict(ct.transform(pd.DataFrame(data,columns=feature_cols)))
    pred = pred[0]
    if pred:
        return render_template("output.html",y="This patient will be readmitted ")
    else:
        return render_template("output.html",y="This patient will not be readmitted")

if __name__ == '__main__':
    app.run(debug = True)

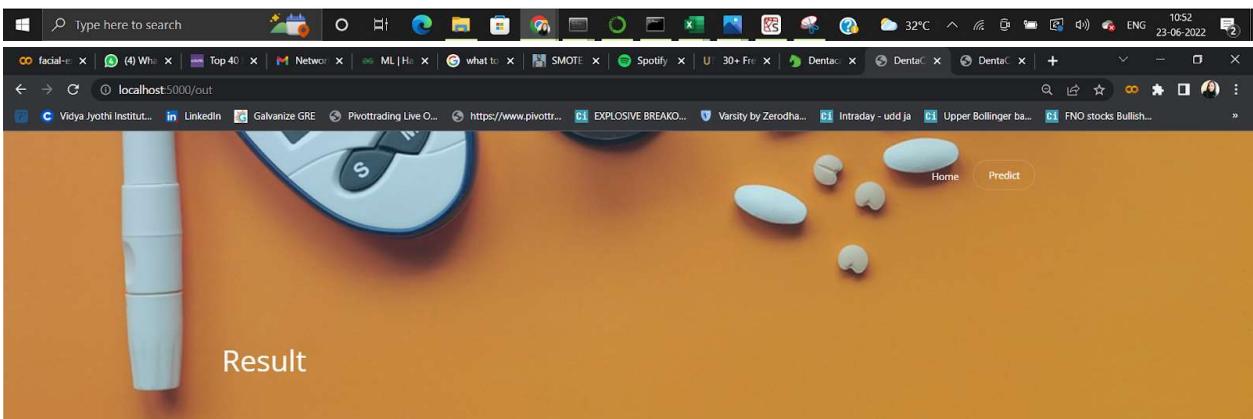
if __name__ == '__main__':
    app.run(debug = True)

```

Output:



This patient will be readmitted



This patient will not be readmitted



